

COMP 6721 Applied Artificial Intelligence (Fall 2023)

Lab Exercise #07: Introduction to Deep Learning

Solutions

PyTorch is a deep learning research platform designed for maximum flexibility and speed.¹ While **scikit-learn** offers user-friendly tools for a wide range of machine learning algorithms, focusing mainly on traditional methods, **PyTorch** caters specifically to deep learning. It provides a dynamic environment that allows for intricate model designs and optimizations. To gain a basic understanding of how to implement an Artificial Neural Network using the **PyTorch** library, in the subsequent questions, you will implement both a simple MLP and a convolutional neural network for a specific image classification task.

Installation. To set up the necessary environment using **conda**, follow the steps below:

1. First, create a new conda environment. This helps in maintaining a clean workspace and avoids conflicts with other packages. Open your terminal or Anaconda prompt and run:

```
conda create --name pytorch_env python=3.8
```

2. Activate the environment:

```
conda activate pytorch_env
```

3. Install **PyTorch** and **torchvision** using the official channel:²

```
conda install pytorch torchvision -c pytorch
```

4. Lastly, ensure you have other required libraries such as **matplotlib** for visualization, if not included in your current environment:

```
conda install matplotlib
```

With these steps completed, you should have a working conda environment named **pytorch_env** with all the necessary libraries installed.

¹See <https://pytorch.org/docs/stable/index.html>

²See <https://pytorch.org/get-started/locally/> for all options

Question 1 Let's use PyTorch to implement a multi-layer perceptron for classifying the CIFAR10 dataset (see Figure 1).³ The `torchvision` package⁴ provides data loaders for common datasets such as Imagenet, CIFAR10, MNIST, etc.

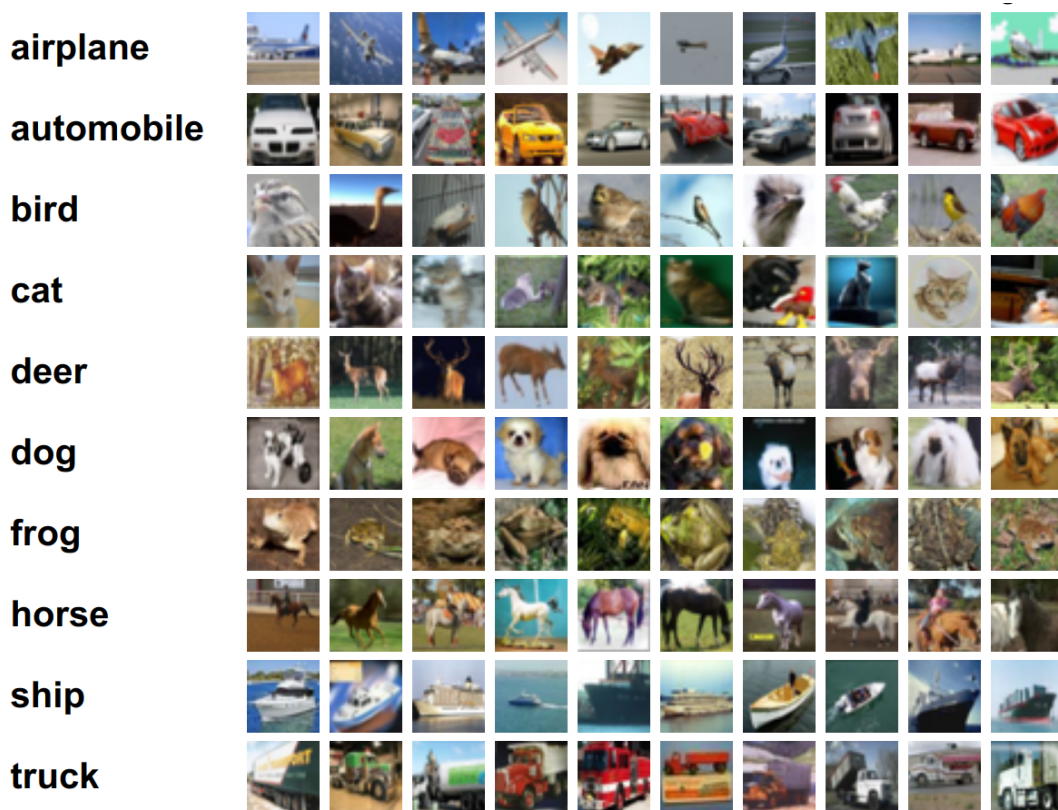


Figure 1: Some example images from the CIFAR-10 dataset

First, utilize the provided code block below, which contains essential Python imports and the `cifar_loader` function, a utility to load the CIFAR-10 dataset. Each CIFAR-10 image is a 32×32 RGB image, giving an input size of $3 \times 32 \times 32 = 3072$. The `cifar_loader` provides train and test data loaders that can be used as iterators. To retrieve the data, standard Python iterators like `enumerate` can be employed. The training dataset uses data augmentation techniques, specifically `RandomHorizontalFlip` and `RandomCrop`, to artificially expand the dataset.⁵ Such augmentations introduce variability, making the model more robust and less prone to overfitting on the training data. After setting the hyper-parameters, where `hidden_size` specifies the hidden dimension and `output_size` represents the output dimension, the dataset is loaded.

³For details on CIFAR10, see <https://en.wikipedia.org/wiki/CIFAR-10>

⁴<https://pytorch.org/docs/stable/torchvision/index.html>

⁵The `RandomHorizontalFlip` augmentation mirrors images, simulating the appearance of objects from different orientations. The `RandomCrop` augmentation helps in focusing on various parts of an image, teaching the model to recognize features irrespective of their position.

```

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.utils.data as td

import torchvision.transforms as transforms
import torchvision.datasets as datasets

# Function to load CIFAR10 dataset
def cifar_loader(batch_size, shuffle_test=False):
    # Normalization values for CIFAR10 dataset
    normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                     std=[0.225, 0.225, 0.225])

    # Loading training dataset with data augmentation techniques
    train_dataset = datasets.CIFAR10('./data', train=True, download=True,
                                     transform=transforms.Compose([
                                         transforms.RandomHorizontalFlip(),
                                         transforms.RandomCrop(32, 4),
                                         transforms.ToTensor(),
                                         normalize
                                     ]))

    # Loading test dataset
    test_dataset = datasets.CIFAR10('./data', train=False,
                                    transform=transforms.Compose([
                                        transforms.ToTensor(),
                                        normalize
                                    ]))

    # Creating data loaders for training and testing
    train_loader = td.DataLoader(train_dataset, batch_size=batch_size,
                                shuffle=True, pin_memory=True)
    test_loader = td.DataLoader(test_dataset, batch_size=batch_size,
                               shuffle=shuffle_test, pin_memory=True)

    return train_loader, test_loader

# Hyperparameters and settings
batch_size = 64
test_batch_size = 64
input_size = 3 * 32 * 32 # 3 channels, 32x32 image size
hidden_size = 50 # Number of hidden units
output_size = 10 # Number of output classes (CIFAR-10 has 10 classes)
num_epochs = 10

```

```
train_loader, _ = cifar_loader(batch_size)
_, test_loader = cifar_loader(test_batch_size)
```

- (a) This is the stage where you'll define the model. You've previously worked with `scikit-learn`'s MLP, so think of this step as specifying the architecture of the neural network. In `PyTorch`, the preferred method to create a neural network is to define a class that inherits from the `nn.Module`⁶ superclass. The `nn.Module` class in `PyTorch` acts as a blueprint that offers functionalities essential for building a variety of deep learning models.

For this exercise, you'll be defining the `MultiLayerFCNet` class, representing a four-layer, fully connected network:

```
model = MultiLayerFCNet(input_size, hidden_size, output_size)
```

This will include an input layer, two hidden layers, and an output layer. Use the hyper-parameter `hidden_size` to specify the number of neurons in each hidden layer. As for the activation functions, employ the ReLU (Rectified Linear Unit) activation for the hidden layers. For the output layer, leverage the `log_softmax` function, which is especially suitable for multi-class classification tasks. The `log_softmax` function provides the logarithm of the softmax values, which, when combined with the negative log likelihood loss, can be used to train models efficiently for classification tasks in `PyTorch`.

Here's a possible solution:

```
class MultiLayerFCNet(nn.Module):
    def __init__(self, D_in, H, D_out):
        # Initialize the parent class, nn.Module
        super(MultiLayerFCNet, self).__init__()

        # Define the input layer (from input dimension to hidden dimension)
        self.linear1 = torch.nn.Linear(D_in, H)

        # Define two hidden layers
        self.linear2 = torch.nn.Linear(H, H)
        self.linear3 = torch.nn.Linear(H, H)

        # Define the output layer (from hidden dimension to output dimension)
        self.linear4 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        # Pass input through the input layer and apply ReLU activation
        x = F.relu(self.linear1(x))

        # Pass through the first hidden layer and apply ReLU activation
```

⁶<https://pytorch.org/docs/stable/generated/torch.nn.Module.html>

```

x = F.relu(self.linear2(x))

# Pass through the second hidden layer and apply ReLU activation
x = F.relu(self.linear3(x))

# Pass through the output layer
x = self.linear4(x)

# Apply log_softmax to get log probabilities for multi-class classification
return F.log_softmax(x, dim=1)

```

The `forward` function overrides the base forward function in `nn.Module` and is required for the `nn.Module` functionality to work correctly. It defines the computational steps for data as it moves through the network.

- (b) Now use PyTorch's `CrossEntropyLoss`⁷ to construct the loss function. This loss is particularly suited for classification problems where the task is to predict one out of multiple classes.

Define an optimizer using `torch.optim`⁸. The optimizer is responsible for updating the weights of our neural network based on the gradients computed during back-propagation. Specifically, we will be using the *Stochastic Gradient Descent (SGD)* optimizer. The learning rate (`lr`) controls the step size during optimization, and `momentum` helps accelerate gradients vectors in the right directions, leading to faster converging.

The first argument passed to the optimizer function contains the parameters we want the optimizer to train. By passing `model.parameters()` to the function, PyTorch effortlessly tracks all the parameters within the model that require training:

```

criterion = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

```

Next, loop over the number of epochs. Within this loop, pass the model outputs and the true labels to the `CrossEntropyLoss` function, defined as `criterion`. Back-propagation is then performed to compute the gradient of the loss with respect to the model parameters. Call `backward()` on the loss variable to execute this. After calculating the gradients using back-propagation, invoke `optimizer.step()` to perform the optimizer's weight update step.

To train the model, we'll iterate over the dataset multiple times, which are known as epochs. For each epoch:

- i. We iterate through batches of data from the `train_loader`.
- ii. Reshape the image data to be suitable for our model.
- iii. Obtain predictions from the model by passing the reshaped images.

⁷<https://pytorch.org/docs/stable/generated/torch.nn.CrossEntropyLoss.html>

⁸<https://pytorch.org/docs/stable/optim.html>

- iv. Compute the loss by comparing the model's predictions to the true labels.
- v. Perform backpropagation to compute gradients.
- vi. Update the model's parameters using the optimizer.
- vii. Track and display the average loss for monitoring purposes.

Here's the code that accomplishes this:

```
# Train the model
for epoch in range(num_epochs):
    avg_loss_epoch = 0 # To keep track of average loss for each epoch
    batch_loss = 0 # Sum of losses for the batches processed
    total_batches = 0 # Total batches processed

    # Iterate over batches of data from train_loader
    for i, (images, labels) in enumerate(train_loader):
        # Reshape images to (batch_size, 32*32*3) - for RGB images of size 32x32
        images = images.reshape(-1, 32 * 32 * 3)

        # Get model predictions for the current batch
        outputs = model(images)

        # Compute the loss between the predicted outputs and true labels
        loss = criterion(outputs, labels)

        # Clear previous gradients
        optimizer.zero_grad()
        # Backpropagate to compute gradients
        loss.backward()
        # Update model parameters
        optimizer.step()

    total_batches += 1
    batch_loss += loss.item()

    # Compute average loss for the current epoch
    avg_loss_epoch = batch_loss / total_batches
    print('Epoch [{}/{}], Average Loss for epoch[{}] = {:.4f}'
          .format(epoch + 1, num_epochs, epoch + 1, avg_loss_epoch))
```

- (c) Finally, we need to monitor the *accuracy* on the test set. To determine the model's predictions, we can use the `torch.max()`⁹ function. This function returns the index of the maximum value in a tensor (an array-like structure). In the context of classification, this index corresponds to the predicted class label. When using `torch.max()`, the first argument should be the output from the model you're examining, and the second argument

⁹<https://pytorch.org/docs/stable/generated/torch.max.html>

specifies the dimension over which to find the maximum value. Your task is to report the accuracy on the test set by comparing the predicted class labels (obtained using `torch.max()`) to the actual labels.

To evaluate the accuracy of the model on the test set:

- i. We initialize two counters: `correct` for the number of correct predictions and `total` for the total number of images.
- ii. Loop through the test data in `test_loader`.
- iii. For each batch of test data, we reshape the images and get predictions from the model.
- iv. We then use `torch.max()` to get the index of the maximum value for each prediction, which represents the predicted class label.
- v. We compare these predicted labels to the true labels to count how many were correct.
- vi. After looping through all test data, we calculate and print the accuracy.

Here's the code that accomplishes this:

```
# Initialize counters
correct = 0
total = 0

# Loop through test data
for images, labels in test_loader:
    # Reshape images to match the input size of the model
    images = images.reshape(-1, 3 * 32 * 32)

    # Get predictions for the current batch of test data
    outputs_test = model(images)

    # Get predicted class labels using torch.max()
    _, predicted = torch.max(outputs_test.data, 1)

    # Update total number of images processed
    total += labels.size(0)

    # Update correct counter by comparing predicted labels to true labels
    correct += (predicted == labels).sum().item()

# Calculate and print accuracy
print('Accuracy of the network on the 10000 test images: %d %%'
      % (100 * correct / total))
```

After running the accuracy calculation code, you might get a result like this:

```
Accuracy of the network on the 10000 test images: 45 %
```

Is this what you'd expect, given our approach?

Considering a dataset with 10 classes, random guessing would yield an accuracy close to 10%. Our model's performance of (here) 45% suggests it has learned patterns from the training data. However, it's essential to recognize the limitations of our current architecture. Simple feed-forward networks aren't optimized for image data. In the next question, we'll explore Convolutional Neural Networks (CNNs) that are specifically tailored for image processing tasks, potentially offering significant improvements in accuracy.

- (d) *Bonus visualization:* If you want to print some random images from the test set with their classification, you can add the code below to your program:

```
# Plot some example images with their predicted labels
import numpy as np
import matplotlib.pyplot as plt

def denormalize(tensor, mean, std):
    for t, m, s in zip(tensor, mean, std):
        t.mul_(s).add_(m)
    return tensor

def plot_results(x=2, y=5):
    # Width per image (inches)
    width_per_image = 2.4

    # Get a batch of images and labels
    data_iter = iter(test_loader)
    images, labels = next(data_iter)

    # Select x*y random images from the batch
    indices = np.random.choice(images.size(0), x*y, replace=False)
    random_images = images[indices]
    random_labels = labels[indices]

    # Get predictions for these images
    random_images_reshaped = random_images.reshape(-1, 3 * 32 * 32)
    outputs = model(random_images_reshaped)
    _, predicted = torch.max(outputs.data, 1)

    # Fetch class names from CIFAR10
    classes = test_loader.dataset.classes

    fig, axes = plt.subplots(x, y, figsize=(y * width_per_image, x *
        ↪ width_per_image))

    # Iterate over the random images and
    # display them along with their predicted labels
```



```

for i, ax in enumerate(axes.ravel()):
    # Denormalize image
    img = denormalize(random_images[i], [0.485, 0.456, 0.406],
        ↳ [0.225, 0.225, 0.225])
    img = img.permute(1, 2, 0).numpy() # Convert image from CxHxW to
        ↳ HxWxC format for plotting

    true_label = classes[random_labels[i]]
    pred_label = classes[predicted[i]]

    ax.imshow(img)
    ax.set_title(f"true='{true_label}', pred='{pred_label}'",
        ↳ fontsize=8)
    ax.axis('off')

plt.tight_layout()
plt.show()

# Call with optional x, y values
plot_results()


```

This will give you an output similar to the one shown in Figure 2.



Figure 2: Some classification results on the CIFAR10 data

Question 2 To improve the performance for image classification, we will use PyTorch to implement more complicated, *deep learning* networks. In this question, you will implement a convolutional neural network (CNN) step-by-step to classify the CIFAR-10 dataset. The CNN architecture that we are going to build can be seen in the diagram below:



- (a) First, use the following code block, which provides the Python imports, the `cifar_loader` to load the dataset, and the hyper-parameters definition:

```

from torch.utils.data import DataLoader
import torch
import torch.nn as nn
import torchvision.transforms as transforms
import torchvision.datasets

num_epochs = 4
num_classes = 10
learning_rate = 0.001

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
train_loader = torch.utils.data.DataLoader(trainset, batch_size=32,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
test_loader = torch.utils.data.DataLoader(testset, batch_size=1000,
                                           shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
  
```

- (b) Now create a class inheriting from the `nn.Module` to define different layers of the network based on provided network architecture above. The first step is to use the `nn.Sequential` module¹⁰ to create sequentially ordered

¹⁰<https://pytorch.org/docs/stable/generated/torch.nn.Sequential.html>

layers in the network. It's a handy way of creating a convolution + ReLU + pooling sequence. In each convolution layer, use **LeakyRelu** for the activation function and **BatchNorm2d**¹¹ to accelerate the training process.

```
class CNN(nn.Module):

    def __init__(self):
        super(CNN, self).__init__()
        self.conv_layer = nn.Sequential(

            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.LeakyReLU(inplace=True),
            nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.LeakyReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(inplace=True),
            nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
        )

        self.fc_layer = nn.Sequential(
            nn.Dropout(p=0.1),
            nn.Linear(8 * 8 * 64, 1000),
            nn.ReLU(inplace=True),
            nn.Linear(1000, 512),
            nn.ReLU(inplace=True),
            nn.Dropout(p=0.1),
            nn.Linear(512, 10)
        )

    def forward(self, x):

        # conv layers
        x = self.conv_layer(x)

        # flatten
        x = x.view(x.size(0), -1)

        # fc layer
```

¹¹<https://pytorch.org/docs/stable/generated/torch.nn.BatchNorm2d.html#torch.nn.BatchNorm2d>

```
x = self.fc_layer(x)
```

```
return x
```

- (c) Before training the model, you have to first create an instance of the **Convolution** class you defined in previous part, then define the loss function and optimizer:
-

```
model = CNN()
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

The following steps are similar to what you've done in previous questions: Loop over the number of epochs and within this loop, pass the model outputs and true labels to the **CrossEntropyLoss** function, defined as **criterion**. Then, perform back-propagation and an optimized training. Call **backward()** on the loss variable to perform the back-propagation. Now that the gradients have been calculated in the back-propagation, call **optimizer.step()** to perform the optimizer training step.

```
total_step = len(train_loader)
```

```
loss_list = []
```

```
acc_list = []
```

```
for epoch in range(num_epochs):
```

```
    for i, (images, labels) in enumerate(train_loader):
```

```
        # Forward pass
```

```
        outputs = model(images)
```

```
        loss = criterion(outputs, labels)
```

```
        loss_list.append(loss.item())
```

```
        # Backprop and optimisation
```

```
        optimizer.zero_grad()
```

```
        loss.backward()
```

```
        optimizer.step()
```

```
        # Train accuracy
```

```
        total = labels.size(0)
```

```
        _, predicted = torch.max(outputs.data, 1)
```

```
        correct = (predicted == labels).sum().item()
```

```
        acc_list.append(correct / total)
```

```
    if (i + 1) % 100 == 0:
```

```
        print('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}, Accuracy: {:.2f}%'  
              .format(epoch + 1, num_epochs, i + 1, total_step, loss_list[-1],  
                      (correct / total) * 100))
```

- (d) Now keep track of the accuracy on the test set. The predictions of the model can be determined by using `torch.max()`.¹²

```
model.eval()
with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in test_loader:
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    print('Test Accuracy of the model on the 10000 test images: {} %'
          .format((correct / total) * 100))
```

- (e) In PyTorch, the learnable parameters (i.e., weights and biases) of a model are contained in the model's parameters. A `state_dict` is simply a Python dictionary object that maps each layer to its parameter tensor. Because `state_dict` objects are Python dictionaries, they can be easily saved, updated, altered, and restored.¹³

Saving the model's `state_dict` with the `torch.save()` function will give you the most flexibility for restoring the model later. To load the models, first initialize the models and optimizers, then load the dictionary locally using `torch.load()`.

Inspect the PyTorch documentation to understand how you can use these functionalities.

To save a model:

```
torch.save(modelA.state_dict(), PATH)
```

To restore a model:

```
modelB = TheModelBClass(*args, **kwargs)
modelB.load_state_dict(torch.load(PATH), strict=False)
```

¹²<https://pytorch.org/docs/stable/generated/torch.max.html>

¹³https://pytorch.org/tutorials/beginner/saving_loading_models.html

Question 3 The goal of `skorch`¹⁴ is to make it possible to use PyTorch with scikit-learn. This is achieved by providing a wrapper around PyTorch that has a scikit-learn interface. Additionally, `skorch` abstracts away the training loop, a simple `net.fit(X, y)` is enough. It also takes advantage of scikit-learn functions, such as `predict`.

- (a) In this section, we will train the same CNN model you developed in the previous question using `skorch` with fewer lines of code.

Let's install `skorch` first:

```
pip install skorch
```

Once the library is installed, we can import the libraries. The next step is to prepare the dataset before training.

```
import matplotlib.pyplot as plt
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.datasets
import torchvision.transforms as transforms
from sklearn.metrics import accuracy_score
from sklearn.metrics import plot_confusion_matrix
from skorch import NeuralNetClassifier
from torch.utils.data import random_split

num_epochs = 4
num_classes = 10
learning_rate = 0.001
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                       download=True, transform=transform)
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)

m = len(trainset)
train_data, val_data = random_split(trainset, [int(m - m * 0.2), int(m * 0.2)])
DEVICE = torch.device("cpu")
y_train = np.array([y for x, y in iter(train_data)])
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

`skorch.NeuralNetClassifier`¹⁵ is a Neural Network class used for classification tasks. Initialize the `NeuralNetClassifier` class then train the CNN

¹⁴<https://skorch.readthedocs.io/en/stable/index.html>

¹⁵<https://skorch.readthedocs.io/en/stable/classifier.html>

model using the method `fit`. Finally print the accuracy score and confusion matrix. Note that CNN is the model you already developed from the previous question using torch.

```
torch.manual_seed(0)
net = NeuralNetClassifier(
    CNN,
    max_epochs=1,
    iterator_train__num_workers=4,
    iterator_valid__num_workers=4,
    lr=1e-3,
    batch_size=64,
    optimizer=optim.Adam,
    criterion=nn.CrossEntropyLoss,
    device=DEVICE
)
net.fit(train_data, y=y_train)
y_pred = net.predict(testset)
y_test = np.array([y for x, y in iter(testset)])
accuracy_score(y_test, y_pred)
plot_confusion_matrix(net, testset, y_test.reshape(-1, 1))
plt.show()
```

- (b) Now let's evaluate the score using K-fold cross-validation. Before the evaluation, we need to make the training set *sliceable* using the class `SliceDataset`, which wraps a torch dataset to make it work with `scikit-learn`. Use the function `cross_val_score(estimator, X, y=None, cv=None)` of the `scikit-learn` library to evaluate the validation accuracy obtained using $K = 5$ folds for K-fold cross-validation.

To achieve this, replace the training process from the previous section with a K-fold cross-validation.

Here is a possible solution:

```
net.fit(train_data, y=y_train)
train_sliceable = SliceDataset(train_data)
scores = cross_val_score(net, train_sliceable, y_train, cv=5,
    scoring="accuracy")
```
